

**Technological Institute of the Philippines**  
938 Aurora Blvd., Cubao, Quezon City

**College of Engineering and Architecture**  
Electronics Engineering Department

Homework II  
**NEURAL STYLE TRANSFER**

**Submitted by:**  
Corcuera, Vincent Cyrus C.  
ECE41S1

**Submitted to:**  
Engr. Christian Lian Paulo P. Rioflorido, MSEE

**October 2022**

We are tasked to optimize an output image to blend 2 images / pictures which is the content image which in this case. The picture of Technological Institute of the Philippines.

I choose 2 artists and their names are **Vincent Van Gogh** and **Benedicto Cabrera**

The first thing we do is to define, import and configure libraries that we need for the project.

```
[ ] import tensorflow_hub as hub
import tensorflow as tf
from matplotlib import pyplot as plt
import numpy as np
import cv2

import os
import tensorflow as tf
# Load compressed models from tensorflow_hub
os.environ['TFHUB_MODEL_LOAD_FORMAT'] = 'COMPRESSED'

import IPython.display as display

import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (12, 12)
mpl.rcParams['axes.grid'] = False

import numpy as np
import PIL.Image
import time
import functools
```

Setup the tensor, download the content image and the style image from the chosen artist. In this case, I use Vincent Van Gogh and Benedicto Cabrera's style

```
[ ] def tensor_to_image(tensor):
    tensor = tensor*255
    tensor = np.array(tensor, dtype=np.uint8)
    if np.ndim(tensor)>3:
        assert tensor.shape[0] == 1
        tensor = tensor[0]
    return PIL.Image.fromarray(tensor)
```

```
content_path = tf.keras.utils.get_file('/content/Technological_Institute_of_the_Philippines_Quezon_City.jpg', 'https://dr
style_path = tf.keras.utils.get_file('/content/download.jpg', 'https://drive.google.com/drive/folders/1_mjFqKw5Z3UhNoyumK
```

Next step will be printing of two image to visualize the output. In order for us to achieve this, we need to define first a function to load each image and limit the dimension of the pixels.

```
▶ def load_img(path_to_img):
    max_dim = 512
    img = tf.io.read_file(path_to_img)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)

    shape = tf.cast(tf.shape(img)[: -1], tf.float32)
    long_dim = max(shape)
    scale = max_dim / long_dim

    new_shape = tf.cast(shape * scale, tf.int32)

    img = tf.image.resize(img, new_shape)
    img = img[tf.newaxis, :]
    return img
```

```
[ ] def imshow(image, title=None):
    if len(image.shape) > 3:
        image = tf.squeeze(image, axis=0)

    plt.imshow(image)
    if title:
        plt.title(title)
```

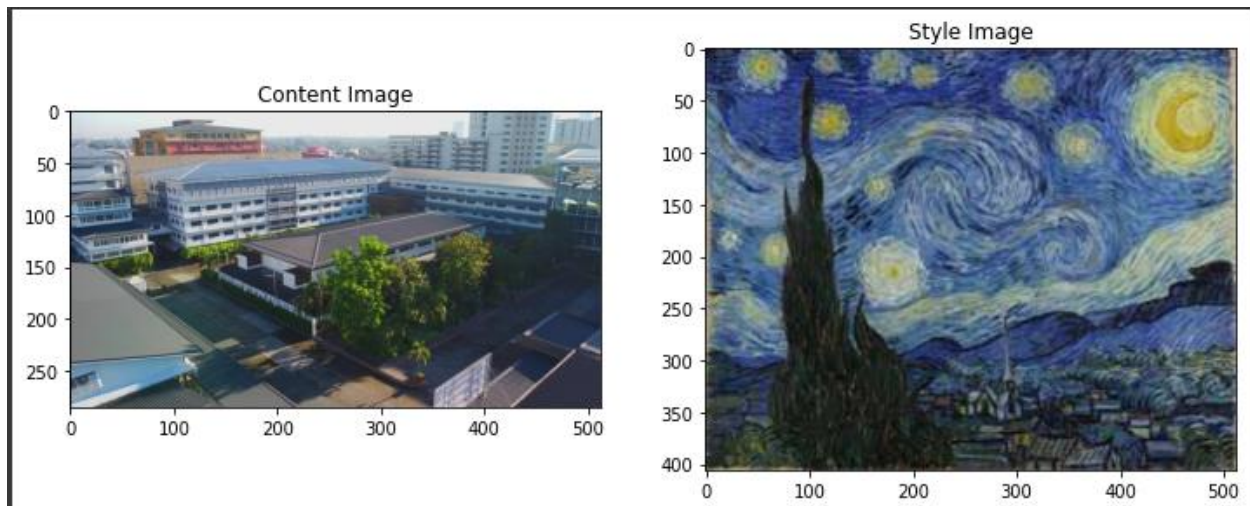
```
[ ] content_image = load_img(content_path)
    style_image = load_img(style_path)

    plt.subplot(1, 2, 1)
    imshow(content_image, 'Content Image')

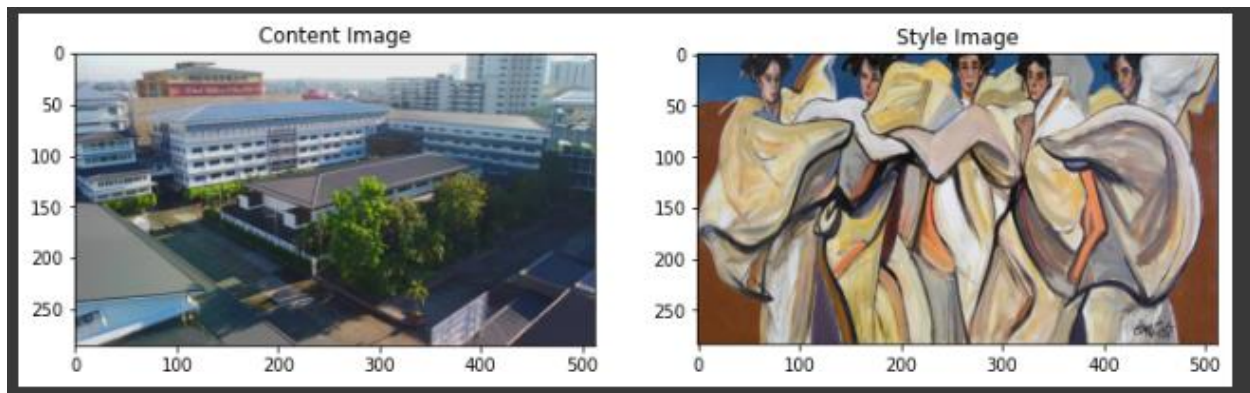
    plt.subplot(1, 2, 2)
    imshow(style_image, 'Style Image')
```

The pictures below are the result of the content that we are doing in the previous steps.

The below pictures are the Technological Institute of the Philippines and Vincent Van Gogh's Starry Night.



The below pictures are the Technological Institute of the Philippines and Ben Cabrera's The Enduring Legacy.



After that, Use VGG19 network architecture to test and run with the images for us to make sure that it is the right picture or image.

```
x = tf.keras.applications.vgg19.preprocess_input(content_image*255)
x = tf.image.resize(x, (224, 224))
vgg = tf.keras.applications.VGG19(include_top=True, weights='imagenet')
prediction_probabilities = vgg(x)
prediction_probabilities.shape
```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19\\_weights\\_tf\\_dim\\_ordering\\_tf\\_574710816/574710816](https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_574710816/574710816) [=====] - 3s 0us/step  
TensorShape([1, 1000])

```
[ ] predicted_top_5 = tf.keras.applications.vgg19.decode_predictions(prediction_probabilities.numpy())[0]
[(class_name, prob) for (number, class_name, prob) in predicted_top_5]
```

Downloading data from [https://storage.googleapis.com/download.tensorflow.org/data/imagenet\\_class\\_index.json](https://storage.googleapis.com/download.tensorflow.org/data/imagenet_class_index.json)  
35363/35363 [=====] - 0s 0us/step  
[('dock', 0.10787788),  
 ('dam', 0.09949343),  
 ('pier', 0.07090867),  
 ('garbage\_truck', 0.046317216),  
 ('crane', 0.038242985)]

```
vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')

print()
for layer in vgg.layers:
    print(layer.name)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19\_weights\_tf\_dim\_ordering\_tf\_80134624/80134624 [=====] - 1s 0us/step
```

input\_2  
block1\_conv1  
block1\_conv2  
block1\_pool  
block2\_conv1  
block2\_conv2  
block2\_pool  
block3\_conv1  
block3\_conv2  
block3\_conv3  
block3\_conv4  
block3\_pool  
block4\_conv1  
block4\_conv2  
block4\_conv3  
block4\_conv4  
block4\_pool  
block5\_conv1  
block5\_conv2  
block5\_conv3  
block5\_conv4  
block5\_pool

```
[ ] content_layers = ['block5_conv2']

style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1']

num_content_layers = len(content_layers)
num_style_layers = len(style_layers)
```

For building the model, use Keras functional API (`tf.keras.applications`). It will extract the intermediate layer values.

We also must specify the inputs since we will use the functional API. That is the reason why we put `model = Model(inputs, outputs)`.

```
[ ] def vgg_layers(layer_names):  
    """ Creates a VGG model that returns a list of intermediate output values."""  
    # Load our model. Load pretrained VGG, trained on ImageNet data  
    vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')  
    vgg.trainable = False  
  
    outputs = [vgg.get_layer(name).output for name in layer_names]  
  
    model = tf.keras.Model([vgg.input], outputs)  
    return model
```

```
▶ style_extractor = vgg_layers(style_layers)  
style_outputs = style_extractor(style_image*255)  
  
#Look at the statistics of each layer's output  
for name, output in zip(style_layers, style_outputs):  
    print(name)  
    print(" shape: ", output.numpy().shape)  
    print(" min: ", output.numpy().min())  
    print(" max: ", output.numpy().max())  
    print(" mean: ", output.numpy().mean())  
    print()
```

This will give us an output of:

```
↳ block1_conv1  
  shape: (1, 406, 512, 64)  
  min: 0.0  
  max: 669.6255  
  mean: 20.270468  
  
block2_conv1  
  shape: (1, 203, 256, 128)  
  min: 0.0  
  max: 2234.299  
  mean: 119.44729  
  
block3_conv1  
  shape: (1, 101, 128, 256)  
  min: 0.0  
  max: 6964.2783  
  mean: 113.765366  
  
block4_conv1  
  shape: (1, 50, 64, 512)  
  min: 0.0  
  max: 14064.77  
  mean: 455.06958  
  
block5_conv1  
  shape: (1, 25, 32, 512)  
  min: 0.0  
  max: 3005.2104  
  mean: 37.969402
```

Use `tf.linalg.einsum` function for our system to implement the Gram matrix calculation concisely.

```
def gram_matrix(input_tensor):
    result = tf.linalg.einsum('bijc,bijd->bcd', input_tensor, input_tensor)
    input_shape = tf.shape(input_tensor)
    num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)
    return result/(num_locations)
```

For extracting the style and content, build a model that returns the tensors of contents and style.

```
class StyleContentModel(tf.keras.models.Model):
    def __init__(self, style_layers, content_layers):
        super(StyleContentModel, self).__init__()
        self.vgg = vgg_layers(style_layers + content_layers)
        self.style_layers = style_layers
        self.content_layers = content_layers
        self.num_style_layers = len(style_layers)
        self.vgg.trainable = False

    def call(self, inputs):
        "Expects float input in [0,1]"
        inputs = inputs*255.0
        preprocessed_input = tf.keras.applications.vgg19.preprocess_input(inputs)
        outputs = self.vgg(preprocessed_input)
        style_outputs, content_outputs = (outputs[:self.num_style_layers],
                                         outputs[self.num_style_layers:])

        style_outputs = [gram_matrix(style_output)
                        for style_output in style_outputs]

        content_dict = {content_name: value
                        for content_name, value
                        in zip(self.content_layers, content_outputs)}

        style_dict = {style_name: value
                     for style_name, value
                     in zip(self.style_layers, style_outputs)}

        return {'content': content_dict, 'style': style_dict}
```

```
extractor = StyleContentModel(style_layers, content_layers)

results = extractor(tf.constant(content_image))

print('Styles:')
for name, output in sorted(results['style'].items()):
    print(" ", name)
    print("   shape: ", output.numpy().shape)
    print("   min: ", output.numpy().min())
    print("   max: ", output.numpy().max())
    print("   mean: ", output.numpy().mean())
    print()

print("Contents:")
for name, output in sorted(results['content'].items()):
    print(" ", name)
    print("   shape: ", output.numpy().shape)
    print("   min: ", output.numpy().min())
    print("   max: ", output.numpy().max())
    print("   mean: ", output.numpy().mean())
```

```

Styles:
  block1_conv1
    shape: (1, 64, 64)
    min: 0.052376863
    max: 15855.283
    mean: 505.8719

  block2_conv1
    shape: (1, 128, 128)
    min: 0.0
    max: 111882.47
    mean: 15729.449

  block3_conv1
    shape: (1, 256, 256)
    min: 0.0
    max: 383264.16
    mean: 16511.375

  block4_conv1
    shape: (1, 512, 512)
    min: 0.0
    max: 5120678.5
    mean: 235370.88

  block5_conv1
    shape: (1, 512, 512)
    min: 0.0
    max: 131512.8
    mean: 1749.0945

```

```

Contents:
  block5_conv2
    shape: (1, 17, 32, 512)
    min: 0.0
    max: 1082.0116
    mean: 16.22038

```

```

[ ] style_targets = extractor(style_image)['style']
    content_targets = extractor(content_image)['content']

[ ] image = tf.Variable(content_image)

[ ] def clip_0_1(image):
    return tf.clip_by_value(image, clip_value_min=0.0, clip_value_max=1.0)

[ ] opt = tf.keras.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)

[ ] style_weight=1e-2
    content_weight=1e4

```

Also, you can see above that we implement style transfer algorithm by running gradient descent. These algorithms will initialize the matching of images.

```

def style_content_loss(outputs):
    style_outputs = outputs['style']
    content_outputs = outputs['content']
    style_loss = tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)
                           for name in style_outputs.keys()])
    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)
                             for name in content_outputs.keys()])
    content_loss *= content_weight / num_content_layers
    loss = style_loss + content_loss
    return loss

```

```

@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))

```

```

train_step(image)
train_step(image)
train_step(image)
tensor_to_image(image)

```



Run a much longer optimization to get the initial output. To clean up the output check for total variation loss and rerun the optimization.

```

import time
start = time.time()

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
        display.clear_output(wait=True)
        display.display(tensor_to_image(image))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))

```





Train step: 1000  
Total time: 4607.4



Train step: 1000  
Total time: 4206.8



```
def high_pass_x_y(image):  
    x_var = image[:, :, 1:, :] - image[:, :, :-1, :]  
    y_var = image[:, 1:, :, :] - image[:, :-1, :, :]  
  
    return x_var, y_var
```

```
[ ] x_deltas, y_deltas = high_pass_x_y(content_image)

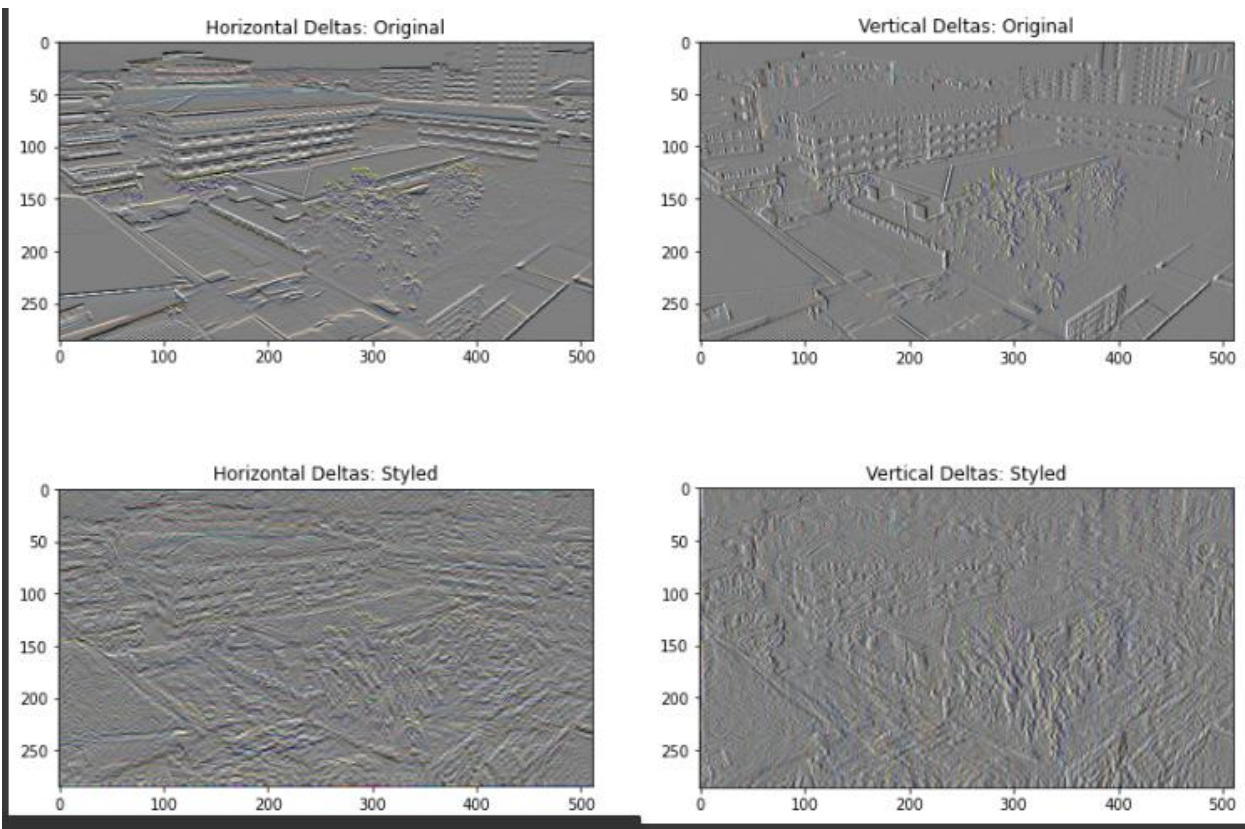
plt.figure(figsize=(14, 10))
plt.subplot(2, 2, 1)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Original")

plt.subplot(2, 2, 2)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Original")

x_deltas, y_deltas = high_pass_x_y(image)

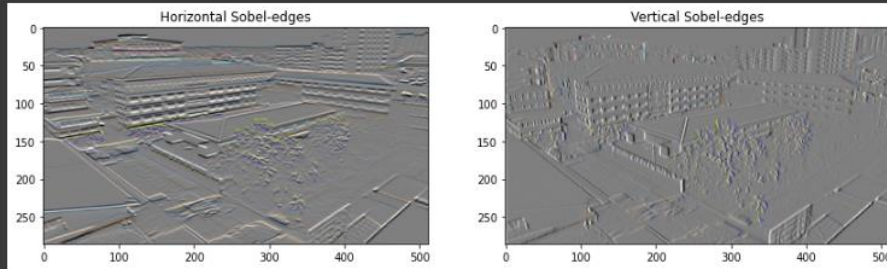
plt.subplot(2, 2, 3)
imshow(clip_0_1(2*y_deltas+0.5), "Horizontal Deltas: Styled")

plt.subplot(2, 2, 4)
imshow(clip_0_1(2*x_deltas+0.5), "Vertical Deltas: Styled")
```



```
[ ] plt.figure(figsize=(14, 10))

sobel = tf.image.sobel_edges(content_image)
plt.subplot(1, 2, 1)
imshow(clip_0_1(sobel[..., 0]/4+0.5), "Horizontal Sobel-edges")
plt.subplot(1, 2, 2)
imshow(clip_0_1(sobel[..., 1]/4+0.5), "Vertical Sobel-edges")
```



```
def total_variation_loss(image):
    x_deltas, y_deltas = high_pass_x_y(image)
    return tf.reduce_sum(tf.abs(x_deltas)) + tf.reduce_sum(tf.abs(y_deltas))
```

```
[ ] total_variation_loss(image).numpy()
```

```
66701.3
```

```
[ ] tf.image.total_variation(image).numpy()
```

```
array([66701.3], dtype=float32)
```

```
[ ] ##Re-run Optimization
total_variation_weight=30
```

```
@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)
        loss += total_variation_weight*tf.image.total_variation(image)

    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))
```

```
[ ] opt = tf.keras.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)
image = tf.Variable(content_image)
```

```

import time
start = time.time()

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
        display.clear_output(wait=True)
        display.display(tensor_to_image(image))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))

```

Lastly, Use this to save output.



Train step: 1000  
Total time: 4634.2

```

file_name = 'VanGogh.png'
tensor_to_image(image).save(file_name)

```



Train step: 400  
.....

```

file_name = 'Bencab.png'
tensor_to_image(image).save(file_name)

```